

Support Vectors Classifier

```
class sklearn.svm.SVC(*, C=1.0, kernel='rbf', degree=3, gamma='scale', coef0=0.0, shrinking=True, probability=False, tol=0.001, cache_size=200,
class_weight=None, verbose=False, max_iter=-1, decision_function_shape='ovr', break_ties=False, random_state=None)

import numpy as np
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler
X = np.array([[ -1, -1], [ -2, -1], [ 1, 1], [ 2, 1]])
y = np.array([1, 1, 2, 2])

from sklearn.svm import SVC
clf = make_pipeline(StandardScaler(), SVC(gamma='auto'))
clf.fit(X, y)

Pipeline(steps=[('standardscaler', StandardScaler()),
                 ('svc', SVC(gamma='auto'))])

print(clf.predict([[ -0.8, -1]]))

[1]
```

Parameters:

- C:float, default=1.0** - Regularization parameter. The strength of the regularization is inversely proportional to C. Must be strictly positive. The penalty is a squared l2 penalty.
- Kernel:{'linear', 'poly', 'rbf', 'sigmoid', 'precomputed'}, default='rbf'** - Specifies the kernel type to be used in the algorithm. It must be one of 'linear', 'poly', 'rbf', 'sigmoid', 'precomputed' or a callable. If none is given, 'rbf' will be used.
- Degree:int, default=3** - Degree of the polynomial kernel function ('poly'). Ignored by all other kernels.
- Gamma:{'scale', 'auto'} or float, default='scale'** - Kernel coefficient for 'rbf', 'poly' and 'sigmoid'.
- Coef0:float, default=0.0** - Independent term in kernel function. It is only significant in 'poly' and 'sigmoid'.
- Shrinking:bool, default=True** - Whether to use the shrinking heuristic.
- Probobaility:bool,default=false** - Whether to enable probability estimates.
- Tol:float, default=1e-3** - Tolerance for stopping criterion.
- cache_size:float, default=200** - Specify the size of the kernel cache (in MB).
- class_weight:dict or 'balanced', default=None** - Set the parameter C of class i to class_weight[i]*C for SVC. If not given, all classes are supposed to have weight one.
- Verbose:bool, default=False** - Enable verbose output.
- max_iter:int, default=-1** - Hard limit on iterations within solver, or -1 for no limit.
- decision_function_shape:{'ovo', 'ovr'}, default='ovr'** - Whether to return a one-vs-rest ('ovr') decision function of shape (n_samples, n_classes) as all other classifiers, or the original one-vs-one ('ovo') decision function of libsvm which has shape (n_samples, n_classes * (n_classes - 1) / 2). However, one-vs-one ('ovo') is always used as multi-class strategy. The parameter is ignored for binary classification.
- break_ties:bool, default=False** - If true, decision_function_shape='ovr', and number of classes > 2, predict will break ties according to the confidence values of decision_function; otherwise the first class among the tied classes is returned.
- random_state:int, RandomState instance or None, default=None** - Controls the pseudo random number generation for shuffling the data for probability estimates. Ignored when probability is False. Pass an int for reproducible output across multiple function calls.

Attributes:

- **class_weight_**: ndarray of shape (n_classes,) - Multipliers of parameter C for each class. Computed based on the class_weight parameter.
- **classes_**: ndarray of shape (n_classes,) - The classes labels.
- **coef_**: ndarray of shape (n_classes * (n_classes - 1) / 2, n_features) - Weights assigned to the features (coefficients in the primal problem). This is only available in the case of a linear kernel.
- **dual_coef_**: ndarray of shape (n_classes - 1, n_SV) - Dual coefficients of the support vector in the decision function multiplied by their targets. For multiclass, coefficient for all 1-vs-1 classifiers.
- **fit_status_**: int - 0 if correctly fitted, 1 otherwise (will raise warning)
- **intercept_**: ndarray of shape (n_classes * (n_classes - 1) / 2,) - Constants in decision function.
- **support_**: ndarray of shape (n_SV) - Indices of support vectors.
- **support_vectors_**: ndarray of shape (n_SV, n_features) - Support vectors.
- **n_support_**: ndarray of shape (n_classes,), dtype=int32 - Number of support vectors for each class.
- **probA_**: ndarray of shape (n_classes * (n_classes - 1) / 2)
- **probB_**: ndarray of shape (n_classes * (n_classes - 1) / 2) - If probability=True, it corresponds to the parameters learned in Platt scaling to produce probability estimates from decision values. If probability=False, it's an empty array.
- **shape_fit_**: tuple of int of shape (n_dimensions_of_X,) - Array dimensions of training vector X.

How sklearn handles SVMs:

- Support vector machines (SVMs) are a set of supervised learning methods used for classification, regression and outliers detection.
- The advantages of support vector machines are:
 1. Effective in high dimensional spaces.
 2. Still effective in cases where number of dimensions is greater than the number of samples.
 3. Uses a subset of training points in the decision function (called support vectors), so it is also memory efficient.
 4. Versatile: different Kernel functions can be specified for the decision function. Common kernels are provided, but it is also possible to specify custom kernels.
- The disadvantages of support vector machines include:
 1. If the number of features is much greater than the number of samples, avoid over-fitting in choosing Kernel functions and regularization term is crucial.
 2. SVMs do not directly provide probability estimates, these are calculated using an expensive five-fold cross-validation.
- The support vector machines in scikit-learn support both dense (`numpy.ndarray` and convertible to that by `numpy.asarray`) and sparse (any `scipy.sparse`) sample vectors as input.
- However, to use an SVM to make predictions for sparse data, it must have been fit on such data.
- For optimal performance, use C-ordered `numpy.ndarray` (dense) or `scipy.sparse.csr_matrix` (sparse) with `dtype=float64`.